

Amazon Elastic Compute Cloud (EC2)

March 2018

What is EC2 ?

- **Virtual Machines / Servers** in Cloud




Amazon
EC2

- **Resizable** compute capacity
- Complete control of your computing resources
- **Reduces the time required** to obtain and boot new server instances to minutes

EC2 Key Concepts...

- **Amazon Machine Image – AMI** –
 - Template containing software configuration (Operating System, Application Server, Applications)
- **Instance** – Launched from AMI
- **Instance Type** – CPU /Memory/Network Performance

<http://aws.amazon.com/ec2/instance-types/>



EC2 Key Concepts

- **Security Group**
 - Virtual Firewall that controls the traffic for one or more instances
 - Inbound Rules
 - Outbound Rules
- **Key Pair**
 - Public key cryptography to encrypt and decrypt login information
 - Public key to encrypt password
 - Private key to decrypt password
 - Public and Private keys are known as key-pair
 - AWS Instance stores Public key
 - If you lose your private key, there is no way to recover it
- **Elastic IP**
 - Static PUBLIC IP

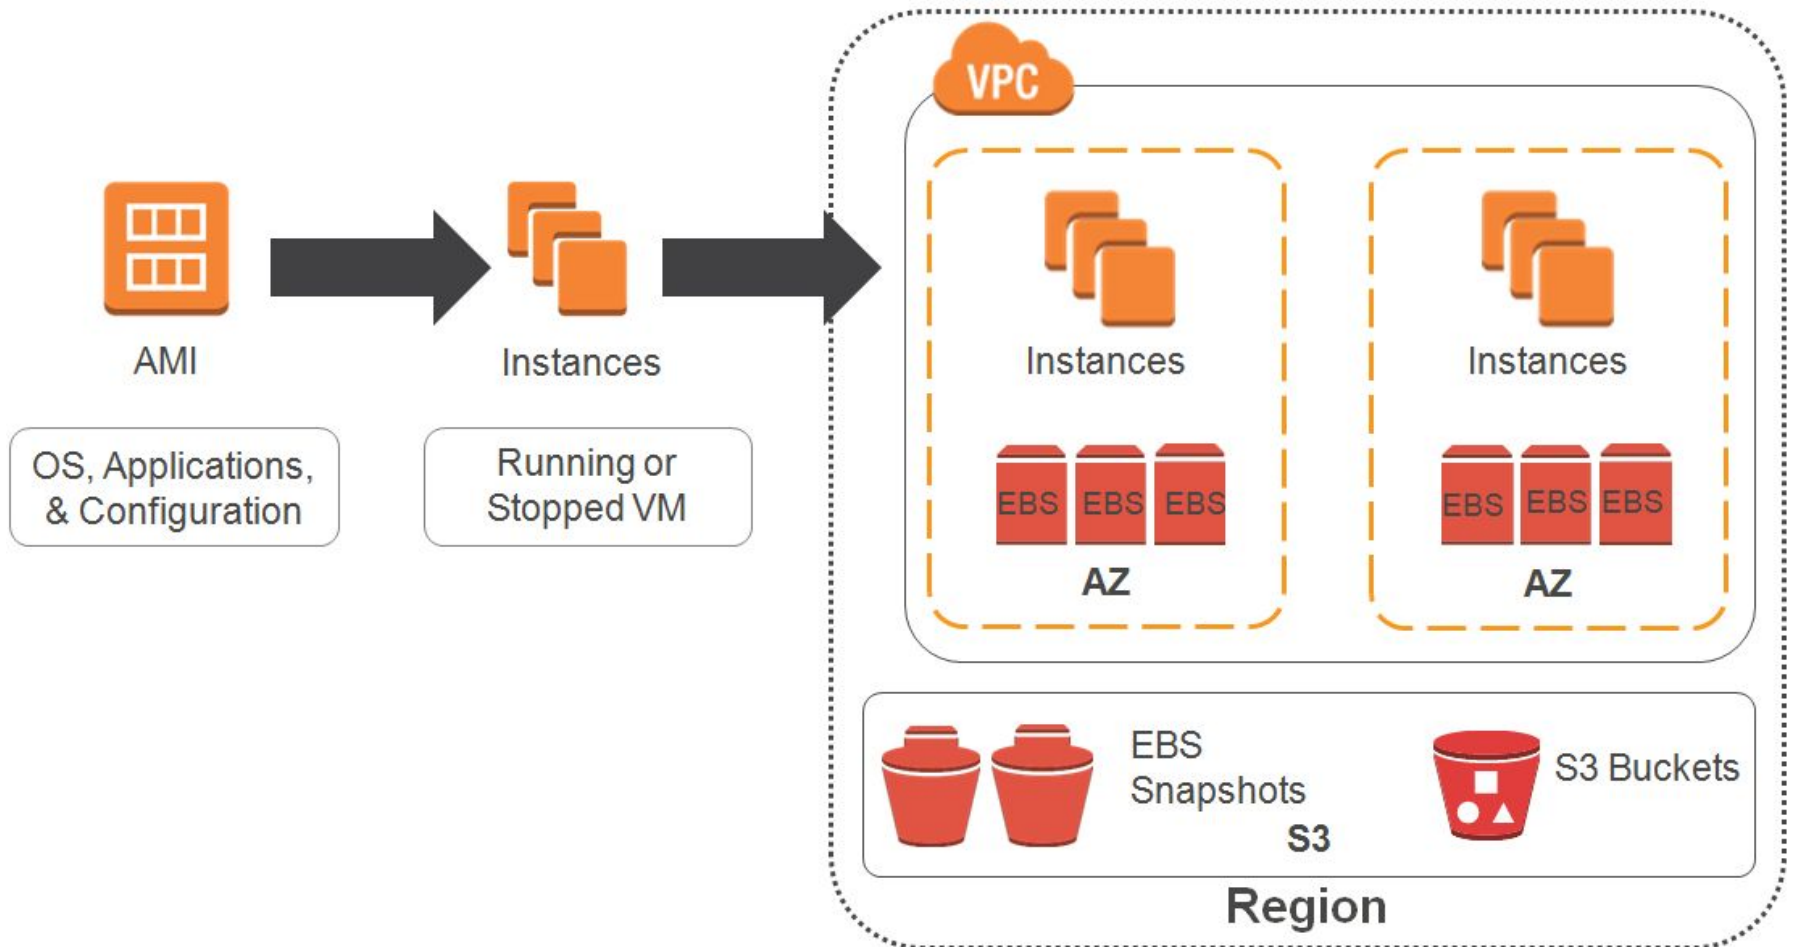
Launching an EC2 Instance

1. Determine the AWS Region in which you want to launch the Amazon EC2 instance.
2. Launch an Amazon EC2 instance from a preconfigured Amazon Machine Image (AMI).
3. Choose an instance type based on CPU, memory, storage, and network requirements.
4. Configure network, IP address, security groups, storage volume, tags, and key pair.

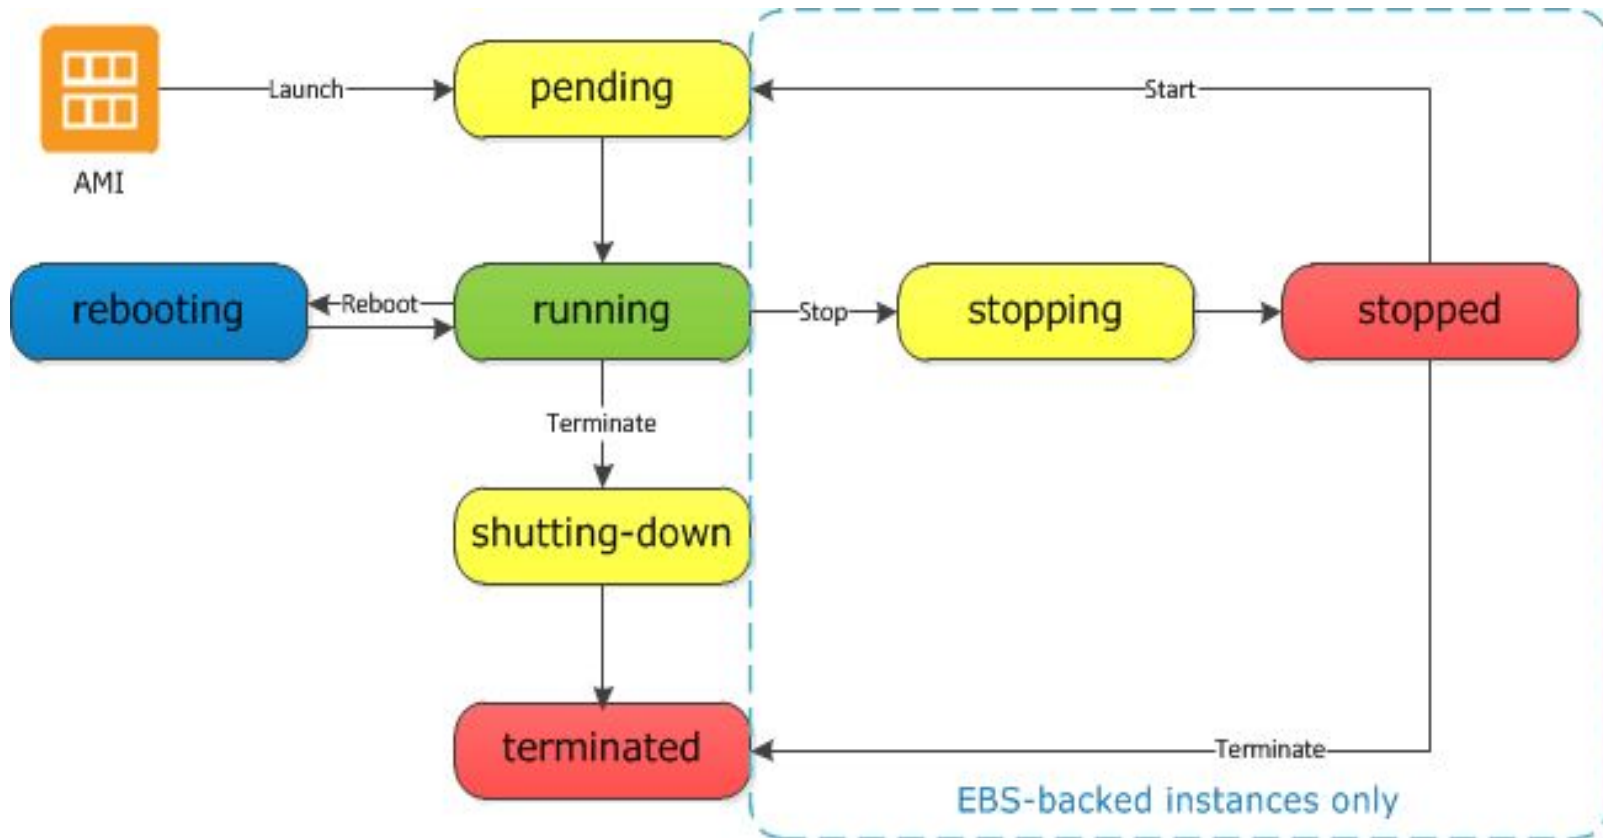
LAB – Launch EC2 Instance

- Login to console
- Create Linux/Windows Instance from AMI
- Finding your instance
- Connect to Instance
- Stop the Instance

Instances in AWS



Instance Lifecycle



Instance Lifecycle States

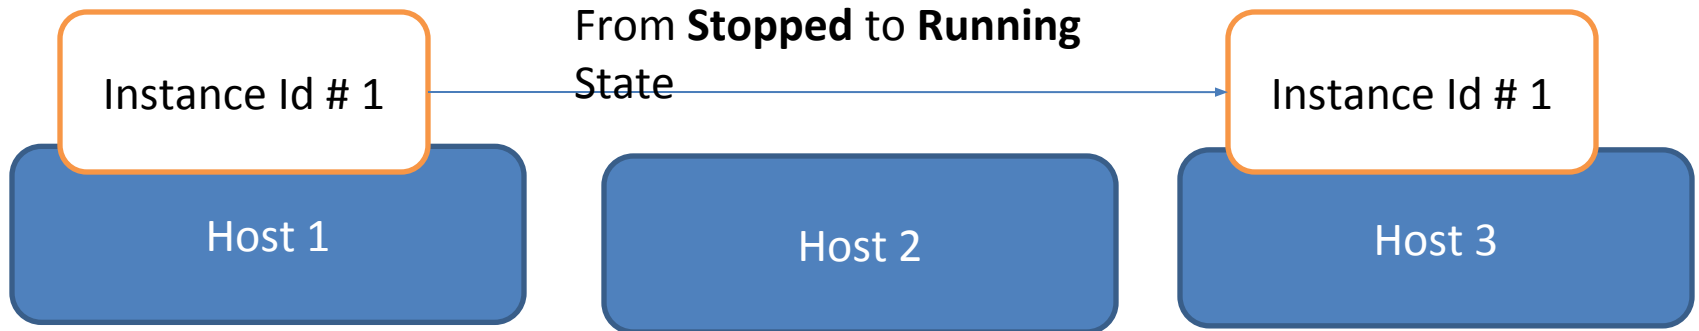
Instance Retirement

- Scheduled to be retired when AWS detects irreparable failure of the underlying hardware hosting the instance
- Either stopped or terminated by AWS

Instance Termination

- Cannot connect to terminated instance
- Cannot recover from terminated instance
- Instance Termination Protection

Instance Changing the host computer



- Loose data on ephemeral storage /instance store volume
- Different PUBLIC IP
- Different private IP in ec2-classic
- Retains same private IP in VPC

Understanding EC2 Instance Details

Win_ser_2012		i-75c22858	t2.micro	us-east-1a	stopped	None	ec2-54-85-32-192
<div>Description</div>							
Instance ID		i-75c22858			Public DNS	ec2-54-85-32-192.compute-1.amazonaws.com	
Instance state		stopped			Public IP	54.85.32.192	
Instance type		t2.micro			Elastic IP	54.85.32.192	
Private DNS		ip-172-31-88-176.ec2.internal			Availability zone	us-east-1a	
Private IPs		172.31.88.176			Security groups	launch-wizard-3, view rules	
Secondary private IPs					Scheduled events	-	
VPC ID		vpc-30d12c51			AMI ID	Windows_Server-2012-R2_RTM-English-64Bit-Base-2014.07.10 (ami-9ade1df2)	
Subnet ID		subnet-e1ce3380			Platform	windows	
Network interfaces		eth0			IAM role	-	
Source/dest. check		True			Key pair name	Win_server_2012_R2_1GB_EBS	
EBS-optimized		False			Owner	-	
Root device type		ebs			Launch time	August 12, 2014 10:06:18 AM UTC+5:30 (1272 hours)	
Root device		-			Termination protection	False	
Block devices		-			Lifecycle	normal	
					Monitoring	basic	
					Alarm status	None	
					Kernel ID	-	
					RAM disk ID	-	
					Placement group	-	
					Virtualization	hvm	
					Reservation	r-31d7c74f	
					AMI launch index	0	
					Tenancy	default	

IP Addresses

- Public IP - 54.85.32.192
- Public DNS - *ec2-54-85-32-192.compute-1.amazonaws.com*
- Elastic IP - 54.85.32.192 – Account Specific
- Private IP - *172.31.88.176*
- Private DNS - *ip-172-31-88-176.ec2.internal*
- Secondary Private IP – Instance can have multiple IP addresses
 - Host multiple websites on a single server
 - Operate network appliance
 - Redirect Internal Traffic to standby instance
 - Management Network

Source/Destination Check

- Indicates whether source/destination checks are performed
- Instance must be the source or destination of any traffic it sends or receives
- Used for **Network Address Translation (NAT)** instance
- Private Subnets communicate to NAT providing its IP
- NAT Instances uses translation and sends the request on behalf of original request

Other Attributes

Owner – 232079927121 – Account Id

Launch Time – The time when instance got STARTED **September 12, 2014 6:28:25 PM UTC+5:30 (529 hours)**

Placement Group - logical grouping of instances within a Single Availability Zone

Monitoring

Alarm Status

Lifecycle - Normal or Spot

Scheduled Events

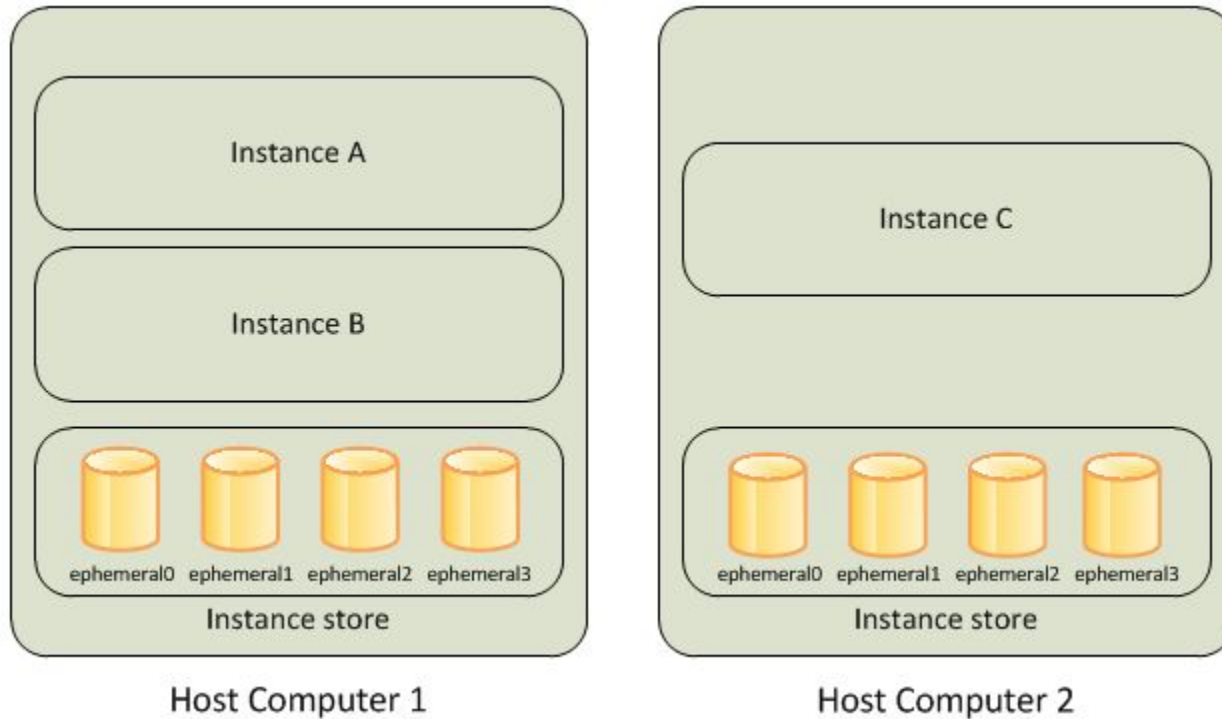
- The number of scheduled events associated with this instance, if applicable
- Types of Scheduled Events
 - Reboot
 - System Maintenance
 - Network
 - Power
 - Instance Retirement
 - Instance Stop

Storage with EC2

Instance Store

- An *instance store* provides temporary block-level storage for your instance
- Storage is located on disks which is physically attached to underlying host
- Ideal for temporary storage of information that changes frequently
 - buffers, caches, scratch data, and other temporary content, or for data that is replicated across a fleet of instances, such as a load-balanced pool of web servers.

Instance Store on Host Computer



Amazon EBS vs. Amazon EC2 Instance Store

Amazon EBS

- Data stored on an Amazon EBS volume can persist independently of the life of the instance.
- Storage is persistent.

Amazon EC2 Instance Store

- Data stored on a local instance store persists only as long as the instance is alive.
- Storage is ephemeral

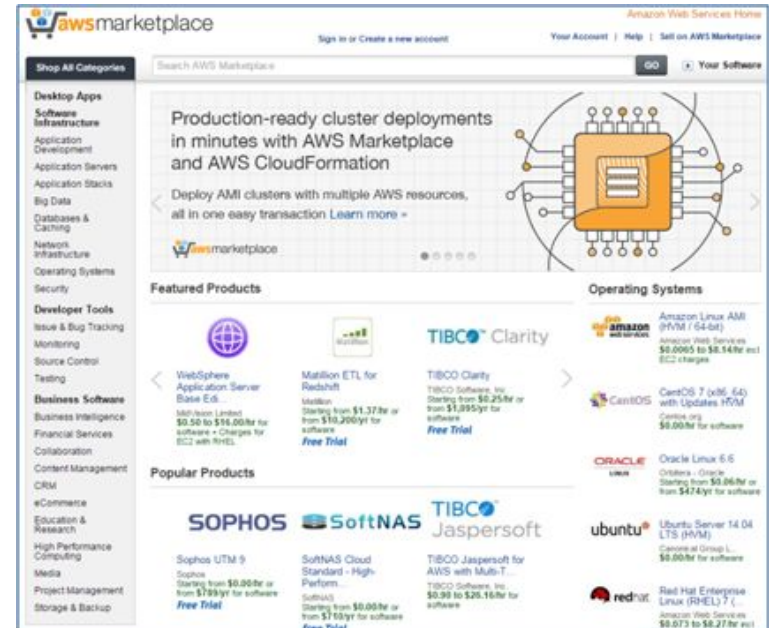
AMI Types - Storage for the Root Device

Boot time	Usually < 1 minute	Usually < 5 minutes
Size limit	16 TiB	10 GiB
Data persistence	The root volume is deleted when the instance terminates. Data on any other Amazon EBS volumes persists after instance termination.	Data on any instance store volumes persists only during the life of the instance.
Charges	Instance usage, Amazon EBS volume usage, and storing your AMI as an Amazon EBS snapshot.	Instance usage and storing your AMI in Amazon S3.
Stopped state	Can be stopped.	Cannot be stopped.

AWS Marketplace - IT Software Optimized for Cloud

Online store to discover, purchase, and deploy IT software on top of the AWS infrastructure.

- Catalog of 2300+ IT software solutions
- Including Paid, BYOL, Open Source, SaaS, & free to try options
- Pre-configured to operate on AWS
- Software checked by AWS for security and operability
- Deploys to AWS environment in minutes
- Flexible, usage-based billing models
- Software charges billed to AWS account
- Includes [AWS Test Drive](#).
- <https://aws.amazon.com/marketplace>



Choosing Right type of Instance

- EC2 instance types are optimized for different use cases and come in multiple sizes. This allows you to optimally scale resources to your workload requirements.
- AWS uses Intel[®] Xeon[®] processors for EC2 instances, providing customers with high performance and value.
- Consider the following when choosing your instances:
 - Core count, memory size, storage size and type, network performance, and CPU technologies.
- **Hurry Up and Go Idle** - A larger compute instance can save you time and money, therefore paying more per hour for a shorter amount of time can be less expensive.

Choosing Instance Family

Instance Family	Some Use Cases
General purpose (t2, m4, m3)	<ul style="list-style-type: none">• Low-traffic websites and web applications• Small databases and mid-size databases
Compute optimized (c4, c3)	<ul style="list-style-type: none">• High performance front-end fleets• Video-encoding
Memory optimized (r3)	<ul style="list-style-type: none">• High performance databases• Distributed memory caches
Storage optimized (i2, d2)	<ul style="list-style-type: none">• Data warehousing• Log or data-processing applications
GPU instances (g2)	<ul style="list-style-type: none">• 3D application streaming• Machine learning

Instance Metadata & User Data

Instance Metadata:

- Is data about your instance.
- Can be used to configure or manage a running instance.

Instance User Data:

- Can be passed to the instance at launch.
- Can be used to perform common automated configuration tasks.
- Runs scripts after the instance starts.

Retrieving Instance Metadata

- To view all categories of instance metadata from within a running instance, use the following URI:

`http://169.254.169.254/latest/meta-data/`

On a Linux instance, you can use:

```
curl http://169.254.169.254/latest/meta-data/  
GET http://169.254.169.254/latest/meta-data/
```

- All metadata is returned as text (content type text/plain)



Adding User Data

- You can specify user data when launching an instance.
- User data can be:
 - Linux script – executed by **cloud-init**
 - Windows batch or PowerShell scripts – executed by **EC2Config** service
- User data scripts run once per **instance-id** by default

User Data Example Linux

- You can specify user data when launching an instance.
- User data can be:
 - Linux script – executed by **cloud-init**
 - Windows batch or PowerShell scripts – executed by **EC2Config** service
- User data scripts run once per **instance-id** by default

User Data Example - Linux

```
#!/bin/sh
```

User data shell scripts must start with the #! characters and the path to the interpreter you want to read the script.

```
yum -y install httpd  
chkconfig httpd on  
/etc/init.d/httpd  
start
```

Install Apache web server
Enable the web server
Start the web server

User Data Example - Windows

```
<powershell>
```

```
Import-Module ServerManager
```



Import the Server Manager module
for Windows PowerShell.

```
Install-WindowsFeature web-server, web-webserver
```

```
Install-WindowsFeature web-mgmt-tools
```

```
</powershell>
```



Install IIS
Install Web Management Tools

Retrieving User Data

To retrieve user data, use the following URI:

`http://169.254.169.254/latest/user-data`

On Linux instance, you can use:

```
curl http://169.254.169.254/latest/user-data/
```

```
$ GET http://169.254.169.254/latest/user-data/
```

Retrieving metadata values

```
[ec2-user@ip-172-31-39-153 scripts]$ curl http://169.254.169.254/latest/meta-data/  
ami-id  
ami-launch-index  
ami-manifest-path  
block-device-mapping/  
hostname  
iam/  
instance-action  
instance-id  
instance-type  
local-hostname  
local-ipv4  
mac  
metrics/  
network/  
placement/  
profile  
public-hostname  
public-ipv4  
public-keys/  
reservation-id  
security-groups  
services/[ec2-user@ip-172-31-39-153 scurl http://169.254.169.254/latest/meta-data/public-ipv4  
13.58.245.251[ec2-user@ip-172-31-39-153 scripts]$
```

Purchasing Options

On-Demand Instances

Pay by the hour.

Reserved Instances

Purchase at significant discount. Instances are always available.

1-year to 3-year terms.

Scheduled Instances

Purchase a 1-year RI for a recurring period of time.

Spot Instances

Highest bidder uses instance at a significant discount.
Spot blocks supported.

Dedicated Hosts

Physical host is fully dedicated to run your instances. Bring your per-socket, per-core, or per-VM software licenses to reduce cost.

Spot Instances

What it is ?

- Bids on spare Amazon EC2 instances
- Price always less than On-Demand EC2 instance
- Prices varies based on demand
- Will not be charged for the interrupted hour
- Types of pot Instance Request
 - One time request
 - Persistent Request
 - request is opened again after your Spot instance is terminated.

How it works ?

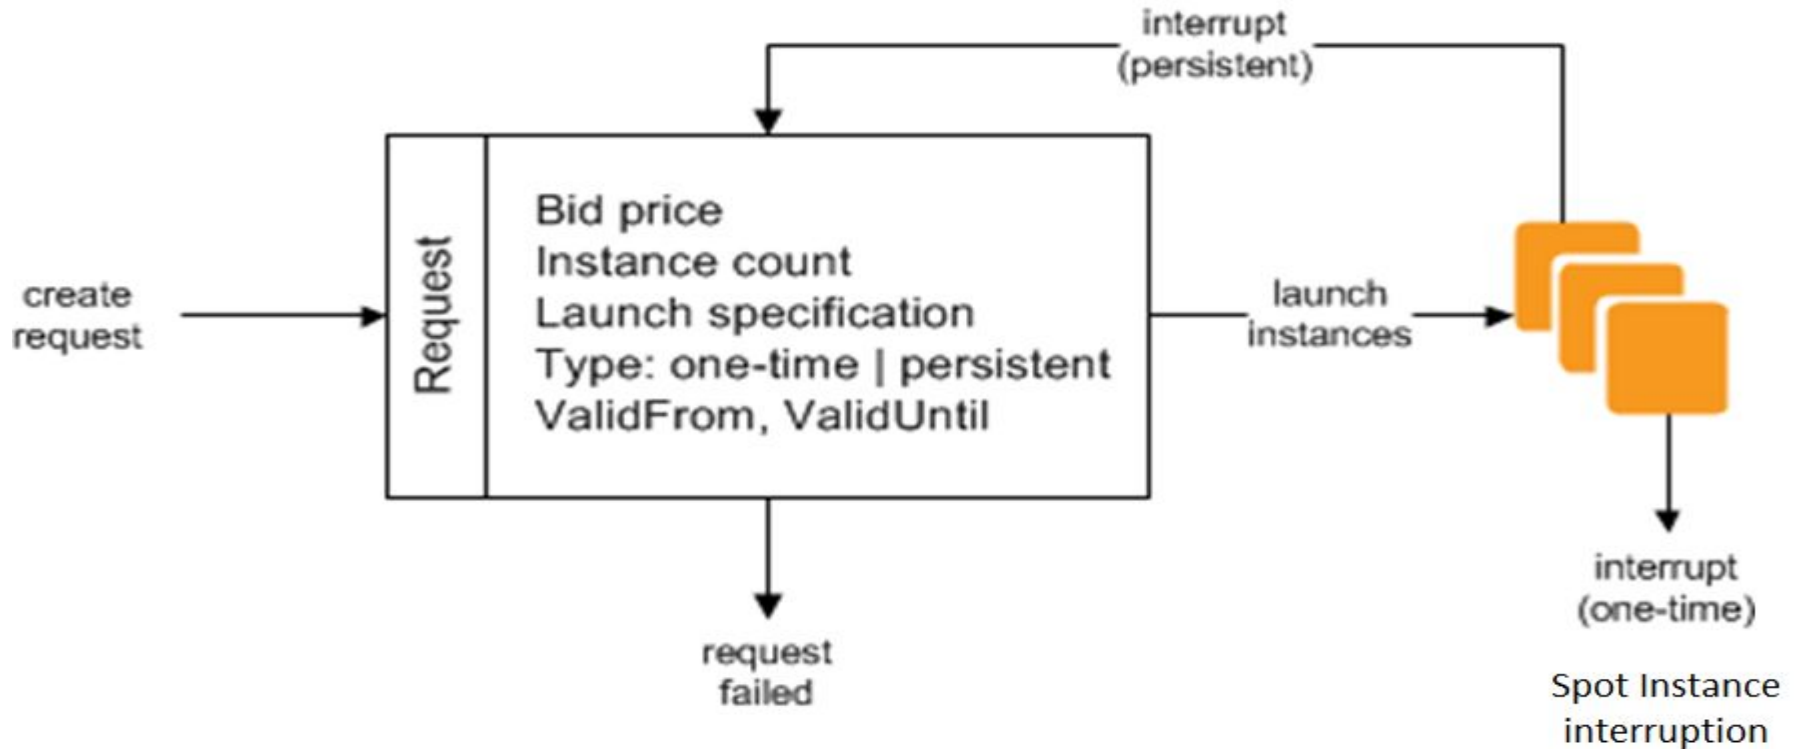
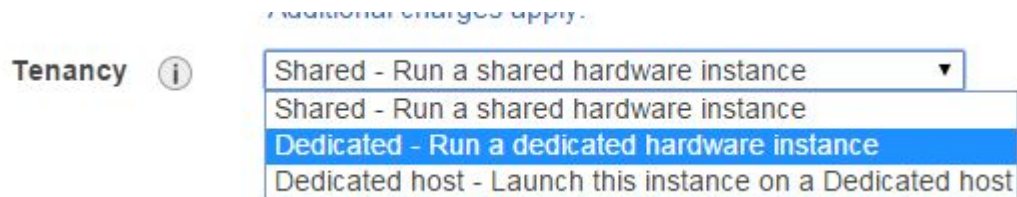


Image Source: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/how-spot-instances-work.html>

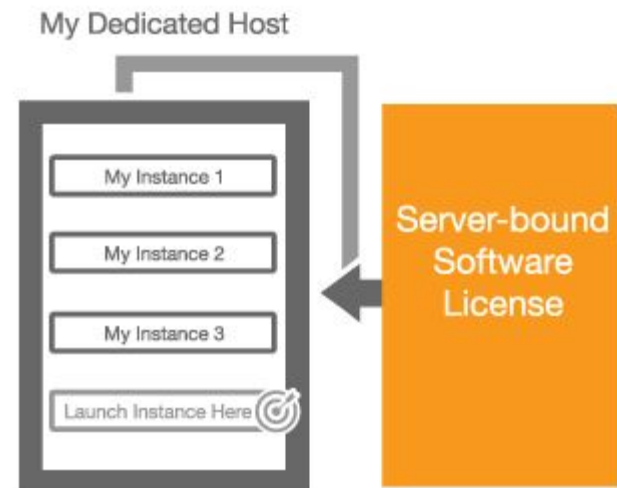
When to Use ?

- Only for time-flexible, interruption-tolerant tasks
- Architect for the potential of interruption
- Large Scale processing
- Payroll processing
- PDF conversion
- Audio - Video Encoding

Dedicated Hosts - New One - Announced on Nov 23 - 2015



- Physical Server with full EC2 instance capacity dedicated to you
- Address compliance requirements
- Allows to reduce costs by allowing to use existing server bound software licences
 - per socket licence
 - per-core licence
 - per-vm licence



Dedicated Hosts vs. Dedicated Instances

Characteristic	Dedicated Instances	Dedicated Hosts
Enables the use of dedicated physical servers	X	X
Per Instance billing (subject to a \$2 per region fee)	X	
Per Host billing		X
Visibility of sockets, cores, host-ID		X
Affinity between a host and instance		X
Targeted instance placement		X
Automatic instance placement	X	X
Add capacity using an allocation request		X

Dedicated Hosts to Instance Mapping

Dedicated Host Attributes			Instance Capacity Per Host by Instance Size						
Instance Family	Sockets	Physical Cores	medium	large	xlarge	2xlarge	4xlarge	8xlarge	10xlarge
c3	2	20	-	16	8	4	2	1	-
c4	2	20	-	16	8	4	2	1	-
g2	2	20	-	-	-	4	-	1	-
m3	2	20	32	16	8	4	-	-	-
d2	2	24	-	-	8	4	2	1	-
r3	2	20	-	16	8	4	2	1	-
m4	2	24	-	22	11	5	2	-	1
i2	2	20	-	-	8	4	2	1	-

- A Dedicated Host is configured to support one instance type at a time.
- e.g. if you allocate a **c3.xlarge** Dedicated Host, you use a Dedicated Host with two sockets and 20 physical cores configured to support up to **8 c3.xlarge** instances.

Pricing

- Price varies by
 - instance family , region and payment option
- Pay hourly
 - for each active Dedicated Host
 - No matter how many instances are launched

General Purpose - Current Generation

m4	\$3.049
m3	\$2.341

Compute Optimized - Current Generation

c4	\$1.939
c3	\$1.848

GPU Instances - Current Generation

g2	\$2.860
----	---------

Memory Optimized - Current Generation

r3	\$3.080
----	---------

Storage Optimized - Current Generation

i2	\$7.502
d2	\$6.072

LAB – Termination Protection

- Login to console
- Enable Termination Protection
- Try to terminate

Amazon Machine Image (AMI)

Amazon Machine Image

- PreConfigured Software Template
- AMI is region specific
- Each Images has a Unique Id

```
login as: ec2-user
Authenticating with public key "imported-openssh-key"

  _ | _ | _ )
  _ | ( _ /   Amazon Linux AMI
  _ | \ _ | _ |

https://aws.amazon.com/amazon-linux-ami/2015.09-release-notes/
3 package(s) needed for security, out of 8 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-10-0-12-35 ~]$ cat /etc/image-id
image_name="amzn-ami-hvm"
image_version="2015.09"
image_arch="x86_64"
image_file="amzn-ami-hvm-2015.09.1.x86_64.ext4.gpt"
image_stamp="1723-6392"
image_date="20151029172510"
recipe_name="amzn ami"
recipe_id="81d8f849-6abb-dda6-8791-35c6-5c72-ce3f-45a14fe7"
[ec2-user@ip-10-0-12-35 ~]$ cat /etc/system-release
Amazon Linux AMI release 2015.09
[ec2-user@ip-10-0-12-35 ~]$ which cloud-init
/usr/bin/cloud-init
[ec2-user@ip-10-0-12-35 ~]$
```

"cloud-init" package

- Application to bootstrap Linux images in a cloud environment
- Open Source
- Developed by Canonical
- Enables to specify actions that should happen to instance at boot time
- Actions can be passed via `user-data` fields
- Use Base Image and pass dynamic data at launch time

For example :

- `action: CONFIG_SSH`
 - Generates host private SSH keys
 - Adds a user's public SSH keys to `.ssh/authorized_keys` for easy login and administration
- `action: PACKAGE_SETUP`
 - Prepares **yum** repo
 - Handles package actions defined in user data

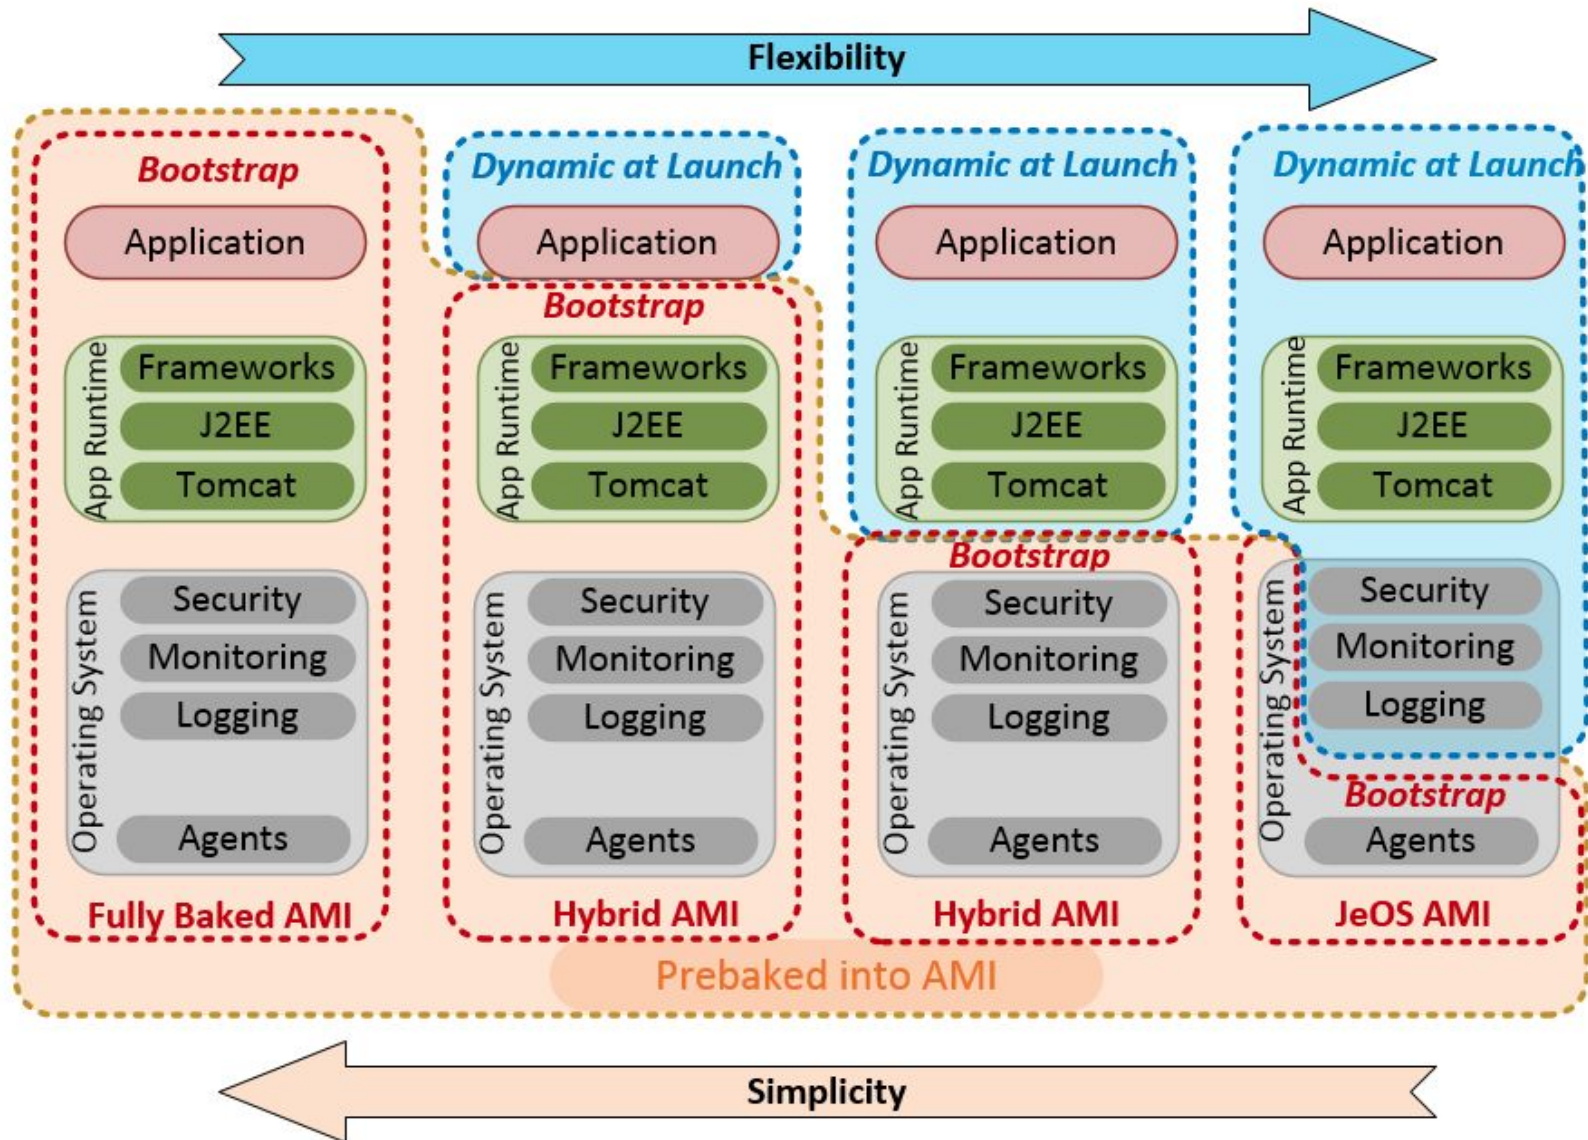
Upgrade the instance

- AMIs are configured to download and install security updates during launch time
- Controlled through 'cloud-init' configuration
- Possible values for repo-upgrade
 - security
 - bugfix
 - all
 - none

```
#cloud-config  
repo_upgrade: security
```

Strategies for Creating Reusable EC2 Instances ? AMIs?

AMI Design Strategies



Choosing Strategy

- How quickly do you need to be able to recover a failing instance or add additional compute capacity?
- Does the workload baseline stay static for a relatively long period?
- Does your server configuration require manual provisioning or configuration?
- Do you need to minimize the complexity of deploying resources to both AWS and on-premises environments?
- Are there existing server provisioning tools or processes that you are trying to align with AWS?

Best Practices

- Avoid embedding passwords, private keys, or other sensitive information in AMIs
- Leverage AWS CloudFormation or a third-party configuration management tool to document and automate AMI creation and updating
- Create a library of reusable, modular templates that can be programmatically assembled to create different types of AMIs
- Instrument AMIs with a standard bootstrapping capability that allows the instance to reference runtime information at launch.
- Develop a consistent strategy for tagging AMIs to allow for the easy organization and identification of images and their contents

LAB – Scaling up/down

- Login to console
- Launch Instance as t1.micro
- Stop it
- Change Instance type to t2.small and Launch
- Connect to the instance

AutoScaling

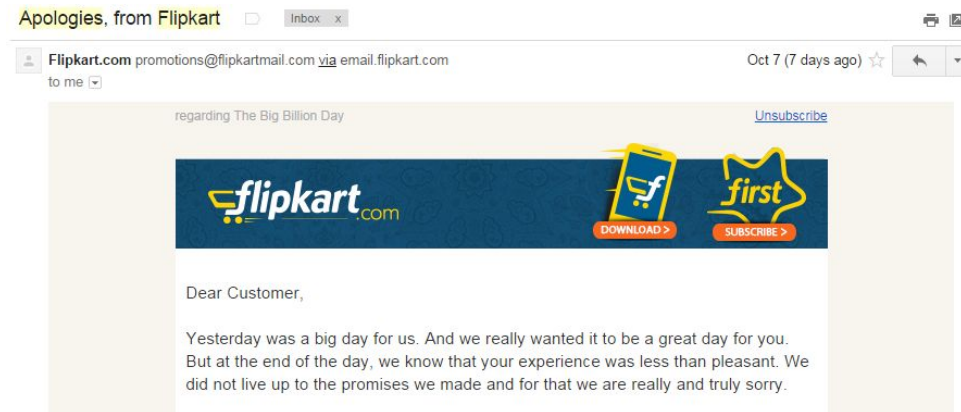
Auto Scaling – Why ?

- Scale Automatically
- High Availability
- Fault Tolerance



Viral Nature

Auto Scaling – Why ?

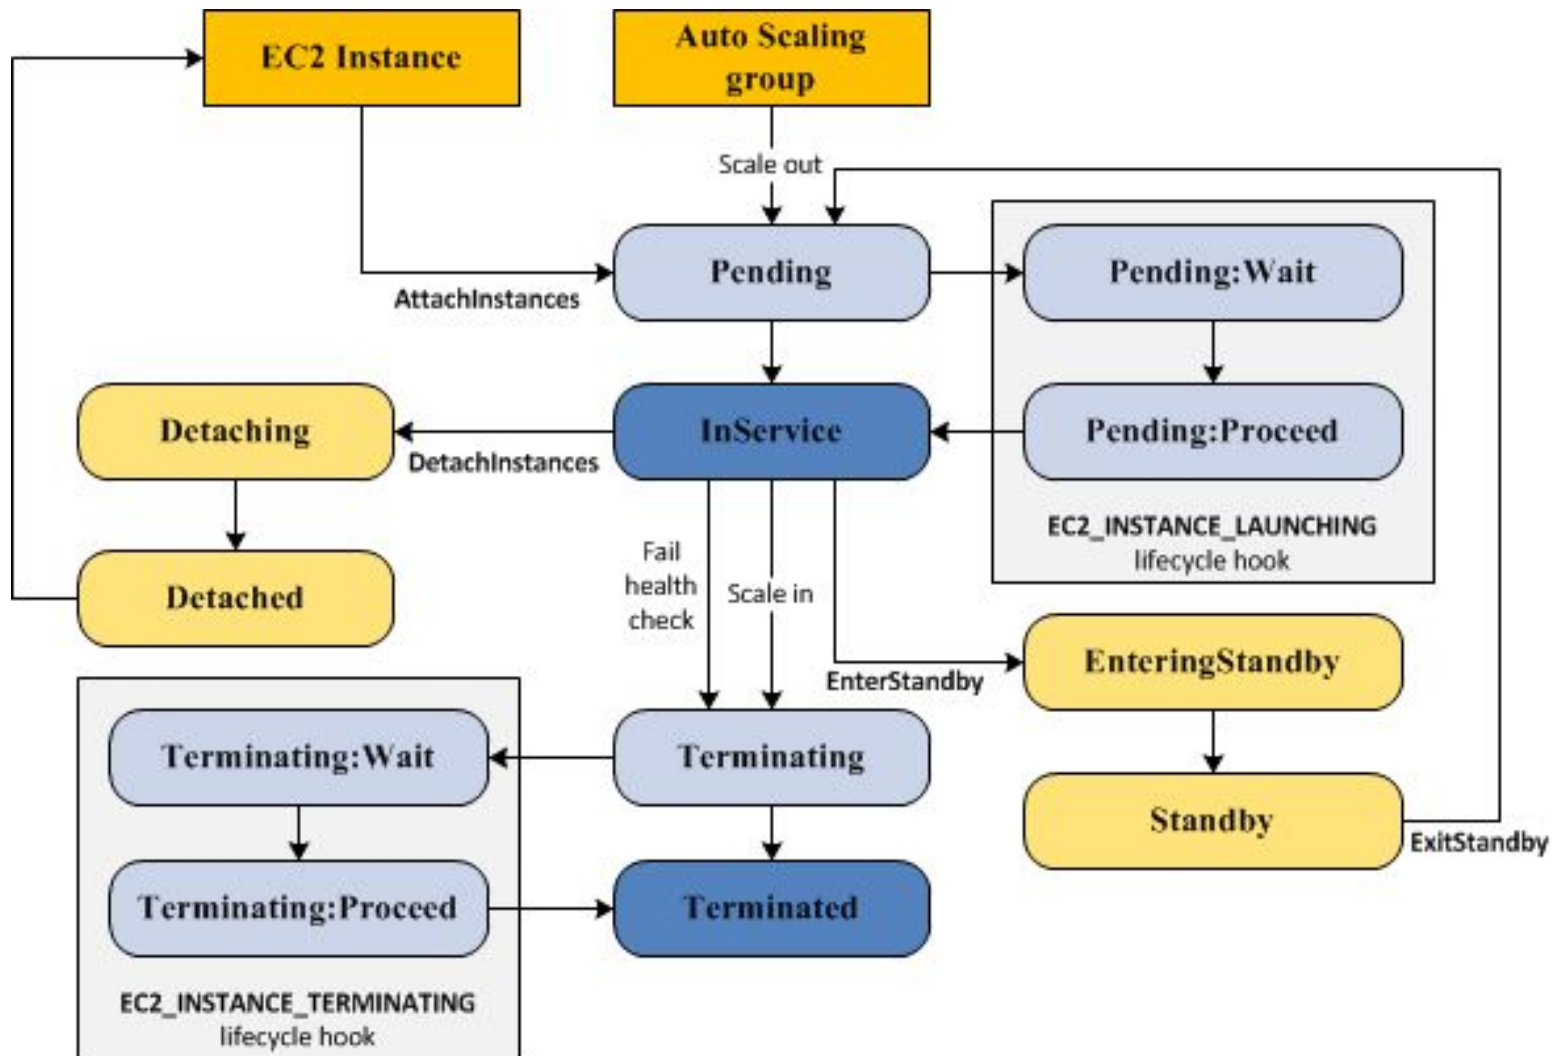


Website Issues We realise that the shopping experience for many of you was frustrating due to errors and unavailability of the website at times. We had deployed nearly 5000 servers and had prepared for 20 times the traffic growth - but the volume of traffic at different times of the day was much higher than this. We are continuing to significantly scale up all our back end systems so that we do a much, much better job next time.

Auto Scaling – Key Concepts

- **Auto Scaling Group**
 - Name
 - Collection of EC2 instances
 - Minimum Instances/ Maximum Instances / Desired Number of Instances
 - Amazon VPC
 - Subnet
 - Metrics and Health Checks
- **Launch Configurations (LC)**
 - Template that will be used by auto scaling group
 - AMI ID, instance type, key pair, security groups, and block device mapping for your instances.
- **Scaling Plan**
 - Tells when and how to scale
 - based on dynamic configuration (cloudwatch alarm)
 - scheduled
- **Integrated with Cloud Watch Alarms**

Auto Scaling – EC2 Instance Lifecycle




Auto Scaling – Steps

1) Create Launch Configuration

- What to Launch - AMI
- Storage
- Security Group
- Monitoring

2) Create Auto Scaling Group

- Name
- Group Size – Desired Capacity
- Scaling Policies
 - When to Scale up and Scale Down



3) Attach Launch Configuration to Auto Scaling Group

End of Module