

# Docker



Feb 2018

# Familiar territory ?

---

- Inconsistencies across environment
- Development is using JRE 8.0 but PROD is still on JRE 7.0
- Missing libraries
- Missing configuration
- Incompatible OS version

**Java - Write Once and Run Everywhere or  
Write Once and Debug Everywhere ?**

---

**Forget about those as I am going to give you  
a bundle / box / which just works**

**I.e. Container (e.g. Docker Container)**

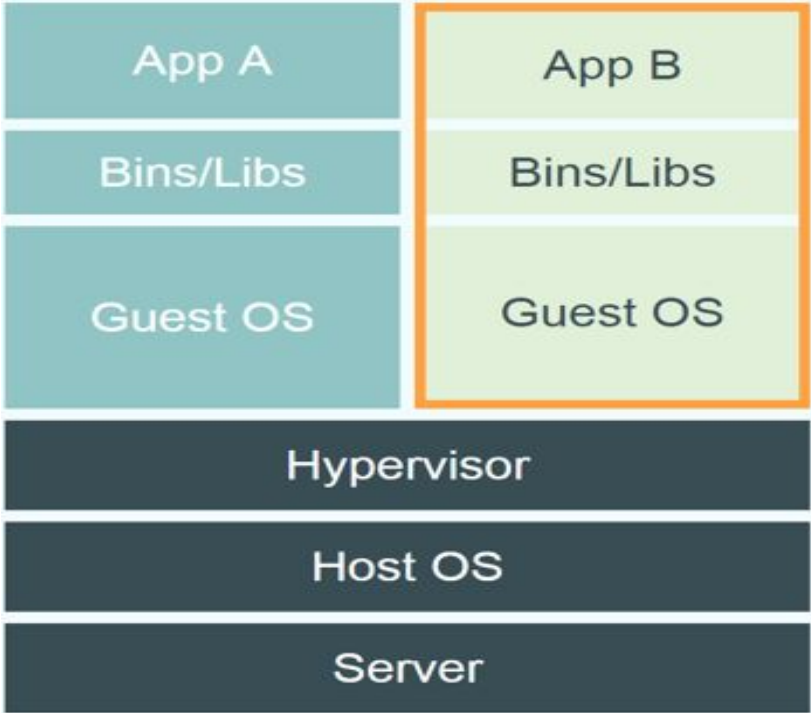
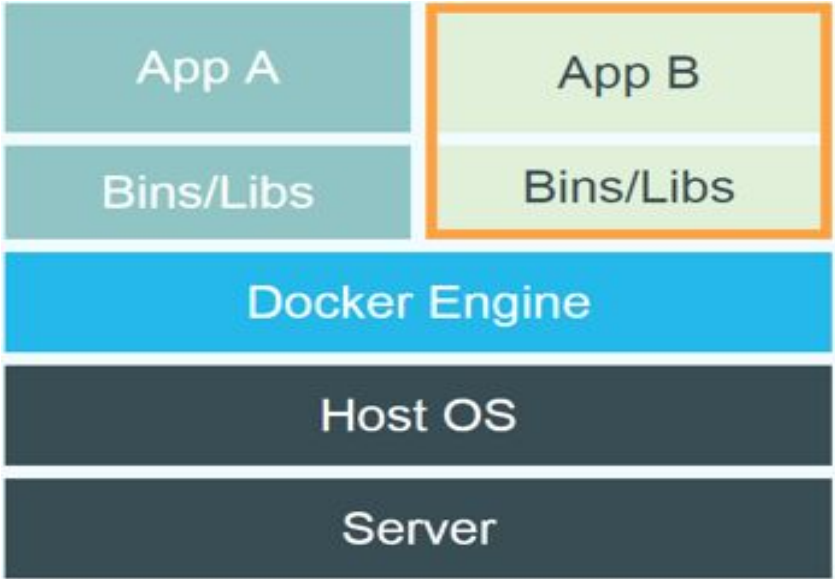
# Overview of Docker

# Docker Platform

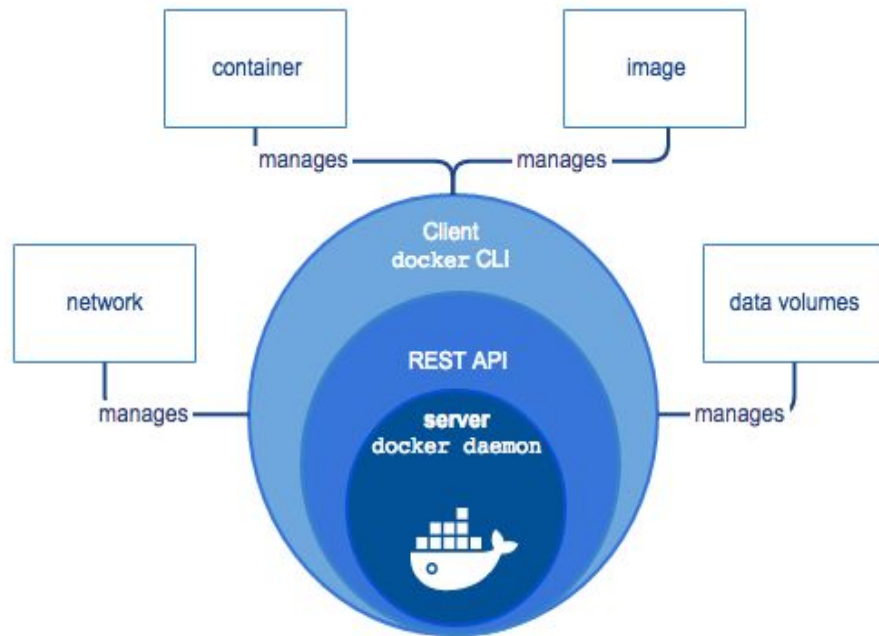
---

- Build, Ship and Run Methodology
  - Provides ability to package and run an application in a loosely isolated environment called a container
  - Isolation and security allow to run many containers simultaneously on a given host
  - Containers are lightweight because they don't need the extra load of a hypervisor
  - Run directly within the host machine's kernel
  - More containers on a given hardware combination than if you were using virtual machines
  - Tooling & Platform to manage lifecycle of containers
  - Works the same whether your production environment is a local data center, a cloud provider, or a hybrid of the two.
-

# Docker Vs Virtual Machine



# Docker Engine



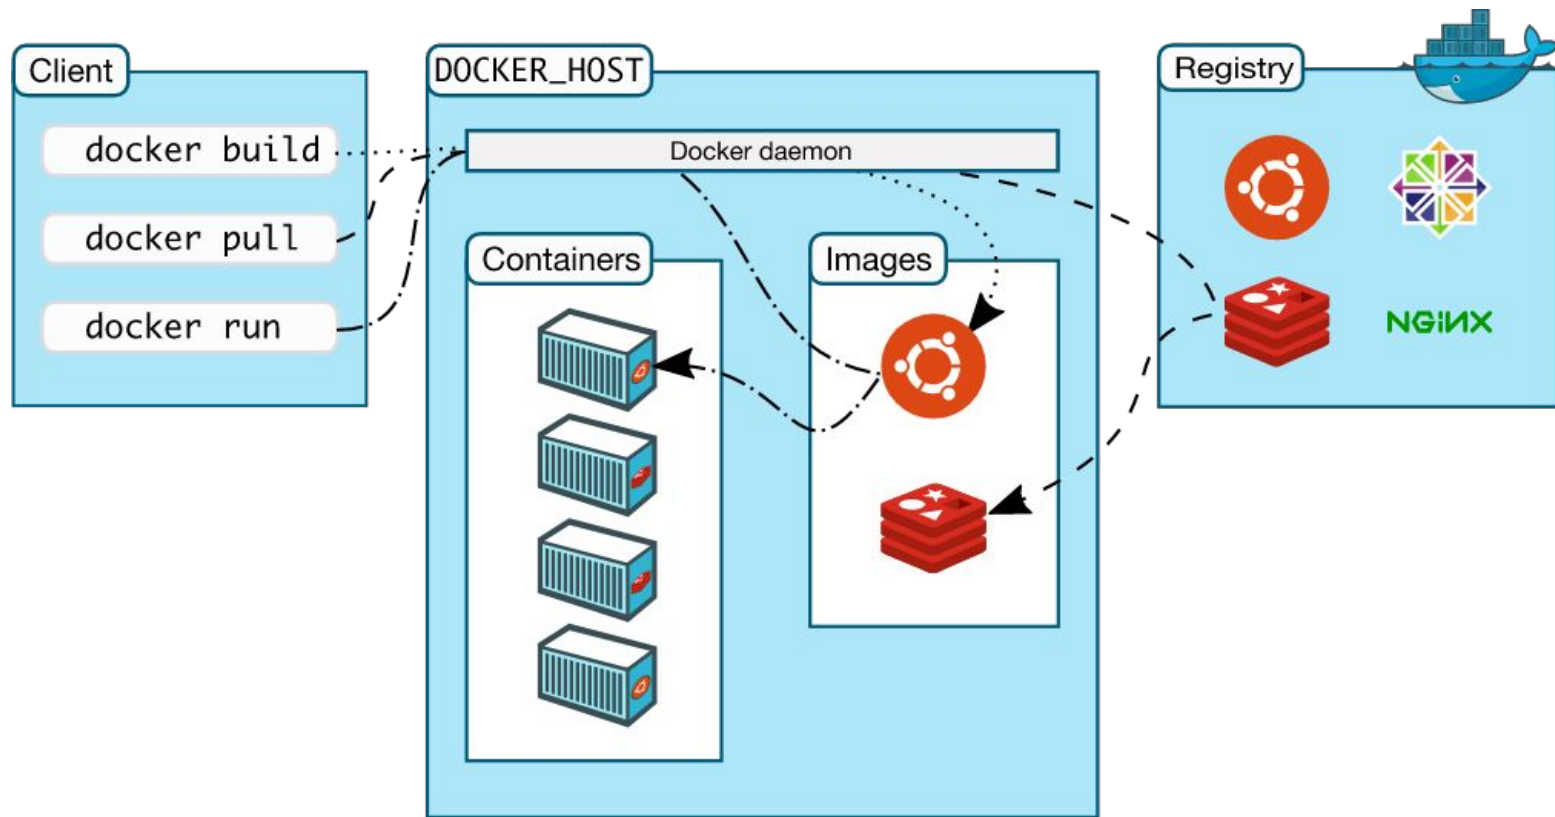
## Client Server application

- Docker daemon (`dockerd`)
- REST API
- Docker CLI

## Docker daemon manages objects such as

- Images
- Containers
- Volumes
- Networks

# Docker Architecture





# Docker Daemon

---

- A.k.a `dockerd`
- listens for Docker API requests and manages Docker objects such as
  - Images
  - Containers
  - Networks &
  - Volumes

# Docker Client

---

- `[docker]` - primary interface for connecting with docker daemon

# Docker Registry

---

- Collection of Repositories
  - Repository is a collection of images
    - Latest
    - V9.0
    - V8.0
    - ...
    - V1.0 and so on
  - Push/Pull images to/from repository
  - Notation
    - `username/repository:tag`
  - Tag - typically used for versions
-

# Images

---

- read-only template with instructions for creating a Docker container
  - *based on* another image, with some additional customization
    - E.g. you may build an image which is based on the **ubuntu** image, but installs the Apache web server and your application, as well as the configuration details needed to make your application run.
  - Image is built using *Dockerfile*
    - Provides simple syntax for defining the steps needed to create the image and run it.
    - Each instruction in a Dockerfile creates a layer in the image.
    - When the image change and rebuilt only those layers which have changed are rebuilt. This makes images lightweight, small, and fast
-

# Container

---

- Runtime instances of portable image
  - Primary runtime unit - defined by image and any runtime parameters
  - Container States - Start , Stop, Remove
  - When a container is removed, any changes to its state that are not stored in persistent storage disappear - *ephemeral storage*
  - Can create image from running container
  - What happens when container is created - for details refer <https://docs.docker.com/engine/docker-overview/#docker-objects>
-

# Services

---

- Allows to scale containers across multiple Docker daemons, which all work together as a *swarm* with multiple *managers* and *workers*
  - Each member of a swarm is a Docker daemon, and the daemons all communicate using the Docker API
  - A service allows you to define the desired state, such as the number of replicas of the service that must be available at any given time
  - By default, the service is load-balanced across all worker nodes
  - To the consumer, the Docker service appears to be a single application.
  - Docker Engine supports swarm mode in Docker 1.12 and higher
-

# **Deploying applications using Docker**

# Container...

---

- Runtime instances of portable image
  - **Image**
    - Contains what goes inside the container
    - Images are created based on Dockerfile
    - **Tag** - identifies the version of image
    - Shared via Registry
  - **Dockerfile**
    - Defines what goes inside the container
    - Series of instructions
    - Access to networking elements, disk drives
    - Port mapping
    - E.g. EXPOSE - exposes the container port
-



# ...Container

---

- Container Commands
    - Running container
    - Running container in detached mode or background
-

# Services

---

- Defines how containers behave in production
  - How many replicas to run
  - Restart behaviour
  - Limits on resources
  - Ports to be used
  - Network to be used
- Scale containers across multiple docker daemons
- By default, the service is load-balanced across all worker nodes.
- To the consumer, the Docker service appears to be a single application
- Uses `docker-compose`

```
services:
  web:
    # replace username/repo:tag with your name and image
    image: username/repo:tag
    deploy:
      replicas: 5
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
      restart_policy:
        condition: on-failure
    ports:
      - "80:80"
    networks:
      - webnet
networks:
  webnet:
```

# ...Services

---

- Changing the replicas results in IN-PLACE deployment

# Stack

---

- Group of interrelated services that
    - share dependencies at the top level &
    - can be orchestrated and scaled together
  - Capable of defining entire application
  - Peer Service
    - Service at the same level
    - No dependency on other service
-

**docker-compose**

# Docker Compose

---

- Tool for defining and running multi-container docker application
  - YAML based
  - Single Command to create and start all the services
  - Start, Stop, Rebuild services
-

# **Docker SWARM**

# Docker Swarm

---

- Dockerized Cluster
    - Multi machine docker containers running as a group
  - Group of Docker Containers
    - Notion of SWARM Manager - Initialize Master node
    - Notion of SWARM Worker
      - Worker nodes which have joined cluster
      - Can Leave SWARM cluster
  - Docker daemon can be SWITCHED ON and OFF with SWARM cluster
-



# SWARM Manager Deployment Strategies

---

- **EMPTIEST\_NODE**
  - Deploy it on the node with least number of containers
- **GLOBAL**
  - Equal distribution of containers across nodes
- Strategies are controlled through `docker-compose` file

# Swarm Worker Node

---

- Just do the work
- Can't authorize other nodes to join cluster

# **Consequences & Benefits of Dockerized Approach**

# Docker Organization Consequence

---

The unit of scale being an individual, portable executable has vast implications.

- CI/CD can push updates to one part of a distributed application
- System dependencies aren't a thing you worry about
- Resource density is increased
- Orchestrating scaling behavior is a matter of spinning up **new executables, not new VM hosts.**

# **Managing Data in Docker**

# How to persist data with Docker ?


---

It is possible to store data within the writable layer of a container, but there are some downsides

- data doesn't persist when that container is no longer running
  - it can be difficult to get the data out of the container if another process needs it
  - is tightly coupled to the host machine where the container is running. You can't easily move the data somewhere else.
  - Writing into a container's writable layer requires a [storage driver](#) to manage the filesystem. The storage driver provides a union filesystem, using the Linux kernel. This extra abstraction reduces performance as compared to using *data volumes*, which write directly to the host filesystem.
-

# 3 ways to persist data

- Volumes
  - Stored in `/var/lib/docker/volumes/`
  - Managed by docker
- bind mounts
  - may be stored *anywhere* on the host system.
- tmpfs volumes
  - stored in the host system's memory only



# Revisit Redis Service Configuration

---

```
redis:
  image: redis
  ports:
    - "6379:6379"
  volumes:
    - "/home/docker/data:/data"
  deploy:
    placement:
      constraints: [node.role == manager]
  command: redis-server --appendonly yes
  networks:
    - webnet
```

Which type ?



# Volume (Data Volume)

---

- Created & Managed by Docker
  - Can be created explicitly or docker can create at the time of Service Creation
  - Different lifecycle than Container
    - Even if the container is removed VOLUME remains
  - Data is stored within a directory on the Docker host
  - volume can be shared across containers as R,RW
  - Can have Explicit NAME or anonymous (docker will give random unique name)
  - Using *volume drivers* data can be stored at remote location (cloud providers)
-

# Bind mounts

---

- Available since early days of docker
  - a file or directory on the *host machine* is mounted into a container
  - The file or directory is referenced by its full path on the host machine
  - are very performant, but they rely on the host machine's filesystem
-

# tmpfs mounts

---

- Only in memory
  - can be used by a container during the lifetime of the container, to store non-persistent state or sensitive information.
    - For instance, internally, swarm services use tmpfs mounts to mount [secrets](#) into a service's containers.
-

# Which one to use ?

---

- Whenever in doubt - use VOLUMES
  - Use case for Volumes
    - Sharing data across containers with R / RW control
    - Storing data on remote host or with cloud provider
    - Backup , restore or migrate data to another host
    - `/var/lib/docker/volumes/<volume-name>`
  - Choosing bind mounts
    - Sharing configuration files from the host machine to containers. This is how Docker provides DNS resolution to containers by default, by mounting `/etc/resolv.conf` from the host machine into each container
-

# Which one to use ?

---

- ...Choosing bind mounts
    - Sharing source code or build artifacts between a development environment on the Docker host and a container
    - you may mount a Maven `target/` directory into a container, and each time you build the Maven project on the Docker host, the container gets access to the rebuilt artifacts
  - Using `tmpfs` mounts
    - Sensitive information
    - Performance reasons - Large amount of data which need not be persisted
-

# Key Notes for Using `'volume'` mount

---

- If you start a container with a volume that does not yet exist, Docker creates the volume for you
  - If you start a container which creates a new volume, as above, and the container has files or directories in the directory to be mounted (such as `/app/`), the directory's contents are copied into the volume. The container then mounts and uses the volume, and other containers which use the volume also have access to the pre-populated content.
  - RO, RW mount flags
  - Data is stored under  
`/var/lib/docker/volumes/<vol_name>/_data/...`
-

# Key Notes for Using `bind mount`

---

- Bind propagation
  - Available only for Linux
- Bind propagation refers to whether or not mounts created within a given bind-mount or named volume can be propagated to replicas of that mount

# Key Notes for Using `tmpfs` mount

---

- Available only for Linux
- Bind propagation refers to whether or not mounts created within a given bind-mount or named volume can be propagated to replicas of that mount



# **Docker Networking**

# Networking Layer

---

- Pluggable architecture for Network using Drivers
  - Bridge
    - Default choice
    - Used when your applications run in standalone containers that need to communicate
    - Restricted to SINGLE HOST
  - Host
    - For standalone containers
    - remove network isolation between the container and the Docker host, and use the host's networking directly
    - but you want other aspects of the container to be isolated
-

# ...Networking Layer

---

- Overlay
    - connect multiple Docker daemons together and enable swarm services to communicate with each other
    - For Multi host configuration
  - Macvlan
    - best choice when dealing with legacy applications that expect to be directly connected to the physical network
  - None
    - Disable all networking (usually works with custom driver)
  - Custom
    -
-

# Working with bridge Network

---

- Default network
    - Not recommended for PROD
    - Can communicate to other containers using only 'docker ips'
  - User defined network
    - Recommended for PROD
    - Can communicate with other container using 'docker ips' +
    - Can communicate with other containers using 'names' . This capability is called **automatic service discovery**.
-

# Working with overlay network

---

- Enables the connectivity between containers running on different docker hosts
- Requires the manager node running in Swarm mode
- Not available to containers started with `docker run` that don't run as part of a swarm mode service
- Create one using

```
Docker network create --driver overlay --subnet 10.0.9.0/24 ol-nw
```

---

# Container and Networks

---

- Container can be part of multiple networks
- Containers can be attached and detached from one network to another

# **Using Dockerfile - building custom image**

# Dockerfile

---

- Provides simple syntax for defining the steps needed to create the image and run it
  - Each instruction in a Dockerfile creates a layer in the image
  - **INSTRUCTION** is not case sensitive but convention is to use CAPITAL letters
  - **The first instruction must be `FROM`** in order to specify the *Base Image* from which you are building.
  - When the image change and rebuilt only those layers which have changed are rebuilt. This makes images lightweight, small, and fast
-



# **Docker Use Cases**

## Use cases

---

- Development Environment
  - Environments for Integration Tests
  - Quick evaluation of software
  - Microservices
  - Multi-Tenancy
  - Unified execution environment (dev,test,prod)
-

# Docker Commands

## Simple Command

---

```
docker run ubuntu echo Hello World
```

`Docker` : tells OS that docker is being used

`run` : creates a container to run specified image

`ubuntu` : image name to be launched inside container

`echo` : command to run inside the container

---

# Flags

---

-d : detached (daemonized)

-t : pseudo terminal inside current

-i : interactive

-f : forcefully

-p : tells docker to map required port of container to host

Multiple flags can be used together for specific requirements

-it , -dt, -idt etc...

---

## Other Parameters

---

--name : name of container

Log : shows the standard output of container

Start : starts the container

Stop : stops the container

rm : remove the container

rmi : remove image

Tag : tags the image

ps : shows running containers

Commit : creates image from container

---

# **Supported Platforms**

# Installation Platforms

---

- Various Linux distributions (Ubuntu, Fedora, RHEL, Centos, openSUSE, ...)
- Cloud (Amazon EC2, Google Compute Engine, Rackspace)
- 64-bit Windows 10 Pro; future release will support more Windows OS versions.



Thanks